



How We Keep Your Data Safe

The handling standard behind every Vision Genesis engagement — whether we're working with your customer list, your invoice history, or your live financials.

Your data is one of your most valuable assets. Handing it to anyone — human or AI — deserves a straight answer about how it's protected. Here's exactly how we handle it.

— Client Isolation

Your data lives in a workspace dedicated solely to you. It's never pooled, blended, or stored alongside another client's information. Every engagement is sealed off — what we learn working with you stays with you.

One client's data never touches another's.

— Data Minimization

We expose only the minimum a task requires. Where identifying details — names, account numbers, customer info — aren't needed for the analysis itself, they're stripped or tokenized before processing and re-attached only on our secure side. The AI sees the shape of the problem, not the identities behind it.

The least exposure needed to get the job done.

— No-Train, No-Reuse AI

When we use frontier AI for the heavy thinking, we operate under enterprise terms — never the consumer apps. Under those terms the providers cannot train on your data, never store it persistently, and can never reuse it for any other client. The models are stateless — your information is processed in the moment and gone, never written to their servers, never carried into anyone else's work.

Processed in the moment. Never stored. Never trains a model.

— Your Data Is Yours

Your data, contacts, and financials are yours start to finish — we're a custodian, never an owner. Access is scoped and revocable, we delete on request, and we never sell or share what you give us.

Custodian, never owner — revocable and deletable anytime.

THE CLEAN LINE

Your data is yours

Your list, prices, customers, and records stay yours — to keep, revoke, or delete. Never sold, never shared.

The platform is ours

The technology and reusable building blocks behind your tools are Vision Genesis IP. What you own outright vs. license for your build is spelled out, in writing, in each engagement.

What we learn, we anonymize

To keep making our tools better for everyone we serve, we may use de-identified, pooled insights — like which automations save owners the most time — never your identifiable data.

IN PRACTICE

Whether it's your **200,000-contact email list**, **twelve months of accounts-receivable history**, or the **financials you need analyzed this quarter** — the same standard applies. Every engagement, every data type, every time.

THE STRAIGHT VERSION

Most heavy data work never touches an AI model at all — it's ordinary, secure processing on our side. Where AI does the thinking, it's on the narrowest slice needed, under enterprise terms that forbid training, reuse, or keeping any copy of your data.